

## COGS118 - HW2

By Gurpreet Dhillon, PID: A14220554

### 1) Decision Boundary

1.1 We are given a classifier that performs classification in  $\mathbb{R}^2$  (the space of data points with 2 features  $(x_1, x_2)$ ) with the following decision rule:

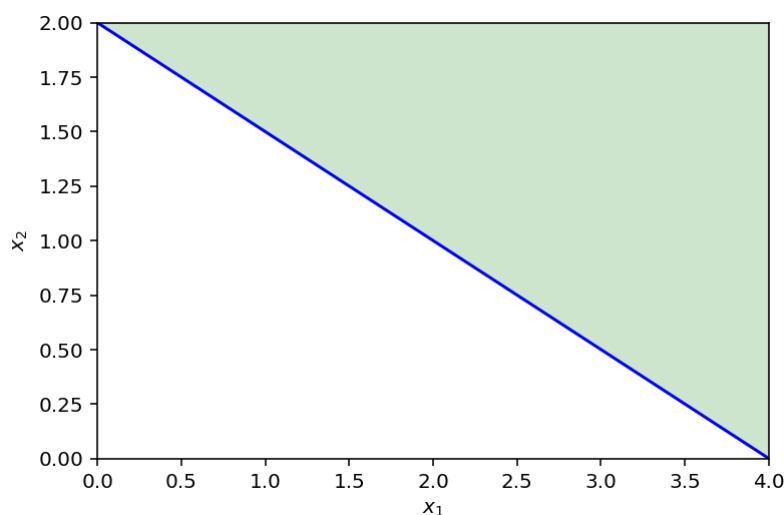
$$h(x_1, x_2) = \begin{cases} 1, & \text{if } x_1 + 2x_2 - 4 \geq 0 \\ 0, & \text{Otherwise} \end{cases}$$

Draw the decision boundary of the classifier and shade the region where the classifier predicts 1. Make sure you have marked the  $x_1$  and  $x_2$  axes and the intercept points on those axes.

```
In [ ]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
```

```
In [42]: x_1 = np.arange(0, 10, 1)
x_2 = 2 - (x_1)/2
plt.plot(x_1, x_2, c='b')
plt.xlabel('$x_1$')
plt.ylabel('$x_2$')
plt.xlim([0,4])
plt.ylim([0,2])
plt.fill_between(x_1, x_2, 2, color= 'g', alpha=0.2)
```

Out[42]: <matplotlib.collections.PolyCollection at 0x1alc74d6d8>



**1.2 We are given a classifier that performs classification in  $\mathbb{R}^2$  (the space of data points with 2 features  $(x_1, x_2)$ ) with the following decision rule:**

$$h(x_1, x_2) = \begin{cases} 1, & \text{if } w_1 x_1 + w_2 x_2 + b \geq 0 \\ 0, & \text{Otherwise} \end{cases}$$

**Here, the normal vector  $w$  of the hyperplane (decision boundary) is normalized, i.e.:**

$$\|w\|^2 = \sqrt{w_1^2 + w_2^2} = 1$$

**1. Compute the parameters  $w_1$ ,  $w_2$  and  $b$  for the decision boundary in Figure 1. Please make sure the prediction of the decision boundary you got is consistent with Figure 1.**

**Hint: Utilize the intercepts in the figure to find the relation between  $w_1$ ,  $w_2$  and  $b$ . Then, substitute it into the normalization constraint to solve the values for the parameters.**

$$\sqrt{\left(\frac{1}{\sqrt{2}}\right)^2 + \left(\frac{1}{\sqrt{2}}\right)^2} = 1$$

$$w_1 = \frac{1}{\sqrt{2}}$$

$$w_2 = \frac{1}{\sqrt{2}}$$

$$w_1 x_1 + w_2 x_2 + b = 0$$

when (0,1):

$$\frac{1}{\sqrt{2}}(1) + b = 0$$

$$b = \frac{-1}{\sqrt{2}}$$

when (1,0):

$$\frac{1}{\sqrt{2}}(1) + b = 0$$

$$b = \frac{-1}{\sqrt{2}}$$

$$h(x_1, x_2) = \begin{cases} 1, & \text{if } \left(\frac{1}{\sqrt{2}}\right)x_1 + \left(\frac{1}{\sqrt{2}}\right)x_2 - \frac{1}{\sqrt{2}} \geq 0 \\ 0, & \text{Otherwise} \end{cases}$$

In [ ]:

2. Compute the predictive labels of the following two data points: A = (3,2), B = (-1,0).

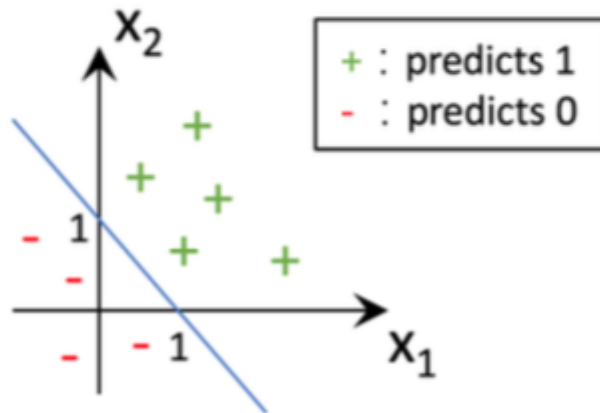


Figure 1: Decision boundary to solve the parameters.

For A:

$$\left(\frac{1}{\sqrt{2}}\right)(3) + \left(\frac{1}{\sqrt{2}}\right)(2) - \frac{1}{\sqrt{2}} \geq 0$$

$$2\sqrt{2} \geq 0 \rightarrow \text{TRUE}$$

A's predicted label is 1 this is because it will fall above the decision boundary.

For B:

$$\left(\frac{1}{\sqrt{2}}\right)(-1) + \left(\frac{1}{\sqrt{2}}\right)(0) - \frac{1}{\sqrt{2}} \geq 0$$

$$-\sqrt{2} \geq 0 \rightarrow \text{FALSE}$$

B's predicted label is 0 this is because it will fall below the decision boundary.

**1.3 We are given a classifier that performs classification on  $R^3$  (the space of data points with 3 features  $(x_1, x_2, x_3)$ ) with the following decision rule:**

$$h(x_1, x_2) = \begin{cases} 1, & \text{if } w_1x_1 + w_2x_2 + w_3x_3 + b \geq 0 \\ 0, & \text{Otherwise} \end{cases}$$

**Here, the normal vector  $w$  of the hyperplane (decision boundary) is normalized, i.e.:**

$$||w||^2 = \sqrt{w_1^2 + w_2^2 + w_3^2} = 1$$

**In addition, we set  $b \leq 0$  to have an unique equation for the decision boundary.**

**1. Compute the parameters  $w_1, w_2, w_3$  and  $b$  for the decision boundary that passes through three points  $A = (3, 2, 4)$ ,  $B = (-1, 0, 2)$ ,  $C = (4, 1, 5)$  in Figure 2.**

**Hint: Note that the normal vector is orthogonal to the hyperplane, which means the normal vector is orthogonal to any vector on the hyperplane. One way to compute the normal vector is to find two nonparallel vectors on the hyperplane, and use the orthogonal property plus the normalization constraint to solve the values for  $w$ . Or you can set the value of  $b$  as any arbitrary value and solve for  $w$ , then come back to solve the true value of  $b$ .**

$$\begin{aligned} B - A &= \{(-1 - 3), (0 - 2), (2 - 4)\} = \{-4, -2, -2\} \\ C - B &= \{(4 + 1), (1 - 0), (5 - 2)\} = \{5, 1, 3\} \\ (B - A) \times (C - B) &= \{-4, 2, 6\} \\ (w_1, w_2, w_3) &= \frac{\{-4, 2, 6\}}{\sqrt{(-4)^2 + (2)^2 + (6)^2}} = \left\{\frac{-4}{56}, \frac{2}{56}, \frac{6}{56}\right\} \\ w_1x_1 + w_2x_2 + w_3x_3 + b &= 0 \\ b &= -(w_1x_1 + w_2x_2 + w_3x_3) \\ B &= (-1, 0, 2) \\ b &= -\left(\left(\frac{-4}{56}\right)(-1) + \left(\frac{2}{56}\right)(0) + \left(\frac{6}{56}\right)(2)\right) \\ b &= \frac{-2}{7} \end{aligned}$$

2. Compute the predictive labels of the following three data points:  $p = (0,0,0)$ ,  $q = (1,0,5)$ .

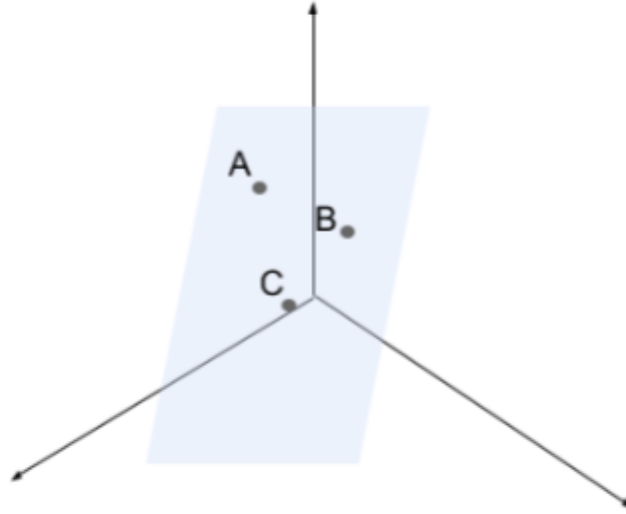


Figure 2: Decision boundary to solve the parameters.

$$\begin{aligned}
 p &= (0, 0, 0) \\
 \left[ \left( \frac{-4}{56} \right)(0) + \left( \frac{2}{56} \right)(0) + \left( \frac{6}{56} \right)(0) \right] + \frac{-2}{7} &\geq 0 \\
 \frac{-2}{7} &\geq 0 \rightarrow \text{FALSE}
 \end{aligned}$$

So,  $P$  is  $< 0$ .  $P$  will be labeled as 0.

$$\begin{aligned}
 q &= (1, 0, 5) \\
 \left[ \left( \frac{-4}{56} \right)(1) + \left( \frac{2}{56} \right)(0) + \left( \frac{6}{56} \right)(5) \right] + \frac{-2}{7} &\geq 0 \\
 \frac{5}{28} &\geq 0 \rightarrow \text{TRUE}
 \end{aligned}$$

So,  $q$  is  $> 0$ .  $q$  will be labeled as 1.

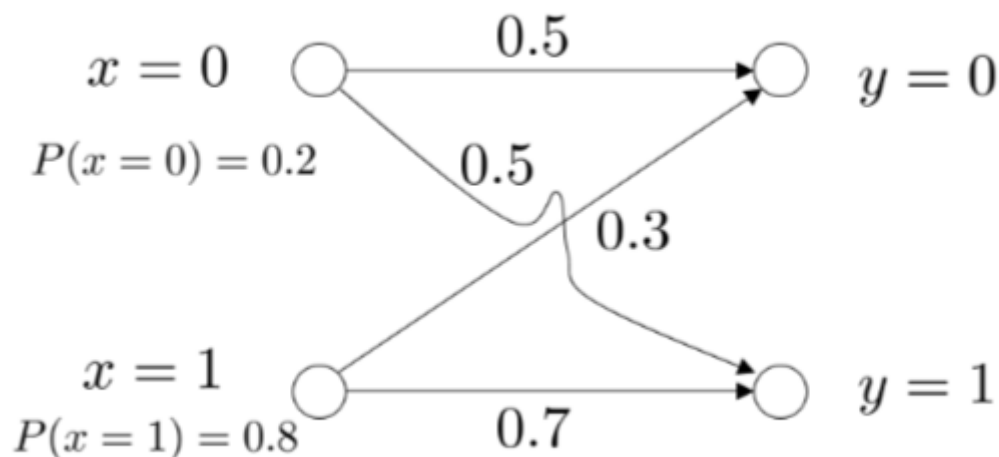
In [ ]:



I dont know why it cuts off here. I dont know why it cuts off here. I dont know why it cuts off here. I dont know why it cuts off here. I dont know why it cuts off here. I dont know why it cuts off here. I dont know why it cuts off here. I dont know why it cuts off here. I dont know why it cuts off here.

### 3) Binary Communication System

For the binary communication system shown below, compute the following probabilities:



The parameter  $p_e$ , probability of error.

The parameter  $1 - p_e$ , probability of no error.

#### a) $P(x=2)$

$x=2$  is not a valid option since because it is not 0 or 1 (We are using a binary coding scheme) So, this is not a valid question.

#### b) $P(y=0|x=1)$

"given that you sent a 1, whats the probability that they recieved a 0?"

$$P(y = 0|x = 1) = .3$$

#### c) $P(y=0)$

$$P(y = 0) = P(y = 0|x = 1)P(x = 1) + P(y = 0|x = 0)P(x = 0) = (.3)(.8) + (.5)(.2) = 0.34$$

#### d) $P(x=1|y=0)$

"given that they recieved a 0, whats the probability that you sent a 1?"

$$P(x = 1|y = 0) = \frac{P(y = 0|x = 1)P(x = 1)}{P(y = 0)} = \frac{(.3)(.8)}{0.34} = 0.705882$$



#### 4) Minimizers and Maximizers

## 4.1 Probability

The joint probability mass function of the random variables  $(x, y)$  is given by the following table. Compute the following:

	$x = 0$	$x = 1$	$x = 2$
$y = 0$	0.2	0.1	0.1
$y = 1$	0.3	0.2	0.1

$$1. (i^*, j^*) = \operatorname{argmax}(i, j) P(x = i, y = j)$$

$$(i^*, j^*) = \operatorname{argmax}(i, j) P(x = i, y = j) = \operatorname{argmax}(x = 2, y = 1) P(x = 2 \cap y = 1) = 0.1$$

[illegible]

**3.  $i^* = \operatorname{argmax}_i P(x=i)$**

$$i^* = \operatorname{argmax}_i P(x = i) = \operatorname{argmax}(2) P(x = 2) = .1 + .1 = 0.2$$

I dont know why it cuts off here. I dont know why it cuts off here. I dont know why it cuts off here. I dont know  
why it cuts off here. I dont know why it cuts off here. I dont know why it cuts off here. I dont know why it cuts off  
here. I dont know why it cuts off here.I dont know why it cuts off here.

## 4.2 Function

(Check Lecture Slide 4, Page44-48) An unknown estimator is given an estimation problem to find the maximizer of the objective function  $G(\theta) \in (0, 2]$ :

$$\theta = \operatorname{argmax}_{\theta} G(\theta)$$

The solution to Eq. 1 by the estimator is  $\theta_1 = 67$ . Given this information, obtain  $\theta^*$  such that

$$\theta^* = \operatorname{argmin}_{\theta} [10 - 3\ln(G(\theta))]$$

$\ln[G(\theta)]$  is an increasing

$-\ln[G(\theta)]$  is not increasing

$-3\ln[G(\theta)]$  is not increasing with scalar 3

$G(\theta)$  is a maximum

so,  $\theta^* = 67$

## 5) Training vs. Testing Errors

```
In [41]: %config InlineBackend.figure_format = 'retina'
import numpy as np
import matplotlib.pyplot as plt
from sklearn import datasets
import pandas as pd
```

```
In [76]: # Iris dataset.
iris = datasets.load_iris()      # Load Iris dataset.

X = iris.data                    # The shape of X is (150, 4), which means
                                # there are 150 data points, each data point
                                # has 4 features.
print(X)
# Here for convenience, we divide the 3 kinds of flowers into 2 groups:
#     Y = 0 (or False): Setosa (original value 0) / Versicolor (original value 1)
#     Y = 1 (or True):  Virginica (original value 2)

# Thus we use (iris.target > 1.5) to divide the targets into 2 groups.
# This line of code will assign:
#     Y[i] = True (which is equivalent to 1) if iris.target[k] > 1.5 (Virginica)
#     Y[i] = False (which is equivalent to 0) if iris.target[k] <= 1.5 (Setosa / Versicolor)

Y = (iris.target > 1.5).reshape(-1,1) # The shape of Y is (150, 1), which means
                                       # there are 150 data points, each data point
                                       # has 1 target value.

X_and_Y = np.hstack((X, Y))        # Stack them together for shuffling.
np.random.seed(1)                  # Set the random seed.
np.random.shuffle(X_and_Y)         # Shuffle the data points in X_and_Y array

print(X.shape)
print(Y.shape)
print(X_and_Y[0])                  # The result should be always: [ 5.8  4.
1.2  0.2  0. ]
```

```
[ [5.1 3.5 1.4 0.2]
[4.9 3. 1.4 0.2]
[4.7 3.2 1.3 0.2]
[4.6 3.1 1.5 0.2]
[5. 3.6 1.4 0.2]
[5.4 3.9 1.7 0.4]
[4.6 3.4 1.4 0.3]
[5. 3.4 1.5 0.2]
[4.4 2.9 1.4 0.2]
[4.9 3.1 1.5 0.1]
[5.4 3.7 1.5 0.2]
[4.8 3.4 1.6 0.2]
[4.8 3. 1.4 0.1]
[4.3 3. 1.1 0.1]
[5.8 4. 1.2 0.2]
[5.7 4.4 1.5 0.4]
[5.4 3.9 1.3 0.4]
[5.1 3.5 1.4 0.3]
[5.7 3.8 1.7 0.3]
[5.1 3.8 1.5 0.3]
[5.4 3.4 1.7 0.2]
[5.1 3.7 1.5 0.4]
[4.6 3.6 1. 0.2]
[5.1 3.3 1.7 0.5]
[4.8 3.4 1.9 0.2]
[5. 3. 1.6 0.2]
[5. 3.4 1.6 0.4]
[5.2 3.5 1.5 0.2]
[5.2 3.4 1.4 0.2]
[4.7 3.2 1.6 0.2]
[4.8 3.1 1.6 0.2]
[5.4 3.4 1.5 0.4]
[5.2 4.1 1.5 0.1]
[5.5 4.2 1.4 0.2]
[4.9 3.1 1.5 0.1]
[5. 3.2 1.2 0.2]
[5.5 3.5 1.3 0.2]
[4.9 3.1 1.5 0.1]
[4.4 3. 1.3 0.2]
[5.1 3.4 1.5 0.2]
[5. 3.5 1.3 0.3]
[4.5 2.3 1.3 0.3]
[4.4 3.2 1.3 0.2]
[5. 3.5 1.6 0.6]
[5.1 3.8 1.9 0.4]
[4.8 3. 1.4 0.3]
[5.1 3.8 1.6 0.2]
[4.6 3.2 1.4 0.2]
[5.3 3.7 1.5 0.2]
[5. 3.3 1.4 0.2]
[7. 3.2 4.7 1.4]
[6.4 3.2 4.5 1.5]
[6.9 3.1 4.9 1.5]
[5.5 2.3 4. 1.3]
[6.5 2.8 4.6 1.5]
[5.7 2.8 4.5 1.3]
[6.3 3.3 4.7 1.6]
```

```
[4.9 2.4 3.3 1. ]
[6.6 2.9 4.6 1.3]
[5.2 2.7 3.9 1.4]
[5.  2.  3.5 1. ]
[5.9 3.  4.2 1.5]
[6.  2.2 4.  1. ]
[6.1 2.9 4.7 1.4]
[5.6 2.9 3.6 1.3]
[6.7 3.1 4.4 1.4]
[5.6 3.  4.5 1.5]
[5.8 2.7 4.1 1. ]
[6.2 2.2 4.5 1.5]
[5.6 2.5 3.9 1.1]
[5.9 3.2 4.8 1.8]
[6.1 2.8 4.  1.3]
[6.3 2.5 4.9 1.5]
[6.1 2.8 4.7 1.2]
[6.4 2.9 4.3 1.3]
[6.6 3.  4.4 1.4]
[6.8 2.8 4.8 1.4]
[6.7 3.  5.  1.7]
[6.  2.9 4.5 1.5]
[5.7 2.6 3.5 1. ]
[5.5 2.4 3.8 1.1]
[5.5 2.4 3.7 1. ]
[5.8 2.7 3.9 1.2]
[6.  2.7 5.1 1.6]
[5.4 3.  4.5 1.5]
[6.  3.4 4.5 1.6]
[6.7 3.1 4.7 1.5]
[6.3 2.3 4.4 1.3]
[5.6 3.  4.1 1.3]
[5.5 2.5 4.  1.3]
[5.5 2.6 4.4 1.2]
[6.1 3.  4.6 1.4]
[5.8 2.6 4.  1.2]
[5.  2.3 3.3 1. ]
[5.6 2.7 4.2 1.3]
[5.7 3.  4.2 1.2]
[5.7 2.9 4.2 1.3]
[6.2 2.9 4.3 1.3]
[5.1 2.5 3.  1.1]
[5.7 2.8 4.1 1.3]
[6.3 3.3 6.  2.5]
[5.8 2.7 5.1 1.9]
[7.1 3.  5.9 2.1]
[6.3 2.9 5.6 1.8]
[6.5 3.  5.8 2.2]
[7.6 3.  6.6 2.1]
[4.9 2.5 4.5 1.7]
[7.3 2.9 6.3 1.8]
[6.7 2.5 5.8 1.8]
[7.2 3.6 6.1 2.5]
[6.5 3.2 5.1 2. ]
[6.4 2.7 5.3 1.9]
[6.8 3.  5.5 2.1]
[5.7 2.5 5.  2. ]
```

```
[5.8 2.8 5.1 2.4]
[6.4 3.2 5.3 2.3]
[6.5 3. 5.5 1.8]
[7.7 3.8 6.7 2.2]
[7.7 2.6 6.9 2.3]
[6. 2.2 5. 1.5]
[6.9 3.2 5.7 2.3]
[5.6 2.8 4.9 2. ]
[7.7 2.8 6.7 2. ]
[6.3 2.7 4.9 1.8]
[6.7 3.3 5.7 2.1]
[7.2 3.2 6. 1.8]
[6.2 2.8 4.8 1.8]
[6.1 3. 4.9 1.8]
[6.4 2.8 5.6 2.1]
[7.2 3. 5.8 1.6]
[7.4 2.8 6.1 1.9]
[7.9 3.8 6.4 2. ]
[6.4 2.8 5.6 2.2]
[6.3 2.8 5.1 1.5]
[6.1 2.6 5.6 1.4]
[7.7 3. 6.1 2.3]
[6.3 3.4 5.6 2.4]
[6.4 3.1 5.5 1.8]
[6. 3. 4.8 1.8]
[6.9 3.1 5.4 2.1]
[6.7 3.1 5.6 2.4]
[6.9 3.1 5.1 2.3]
[5.8 2.7 5.1 1.9]
[6.8 3.2 5.9 2.3]
[6.7 3.3 5.7 2.5]
[6.7 3. 5.2 2.3]
[6.3 2.5 5. 1.9]
[6.5 3. 5.2 2. ]
[6.2 3.4 5.4 2.3]
[5.9 3. 5.1 1.8]]
(150, 4)
(150, 1)
[5.8 4. 1.2 0.2 0. ]
```

In [77]: *# Divide the data points into training set and test set.*

```
X_shuffled = X_and_Y[:, :4]
Y_shuffled = X_and_Y[:, 4]

X_train = X_shuffled[:100] # Shape: (100,4)
Y_train = Y_shuffled[:100] # Shape: (100,)
X_test = X_shuffled[100:] # Shape: (50,4)
Y_test = Y_shuffled[100:] # Shape: (50,)
print(X_train.shape)
print(Y_train.shape)
print(X_test.shape)
print(Y_test.shape)
```

(100, 4)

(100,)

(50, 4)

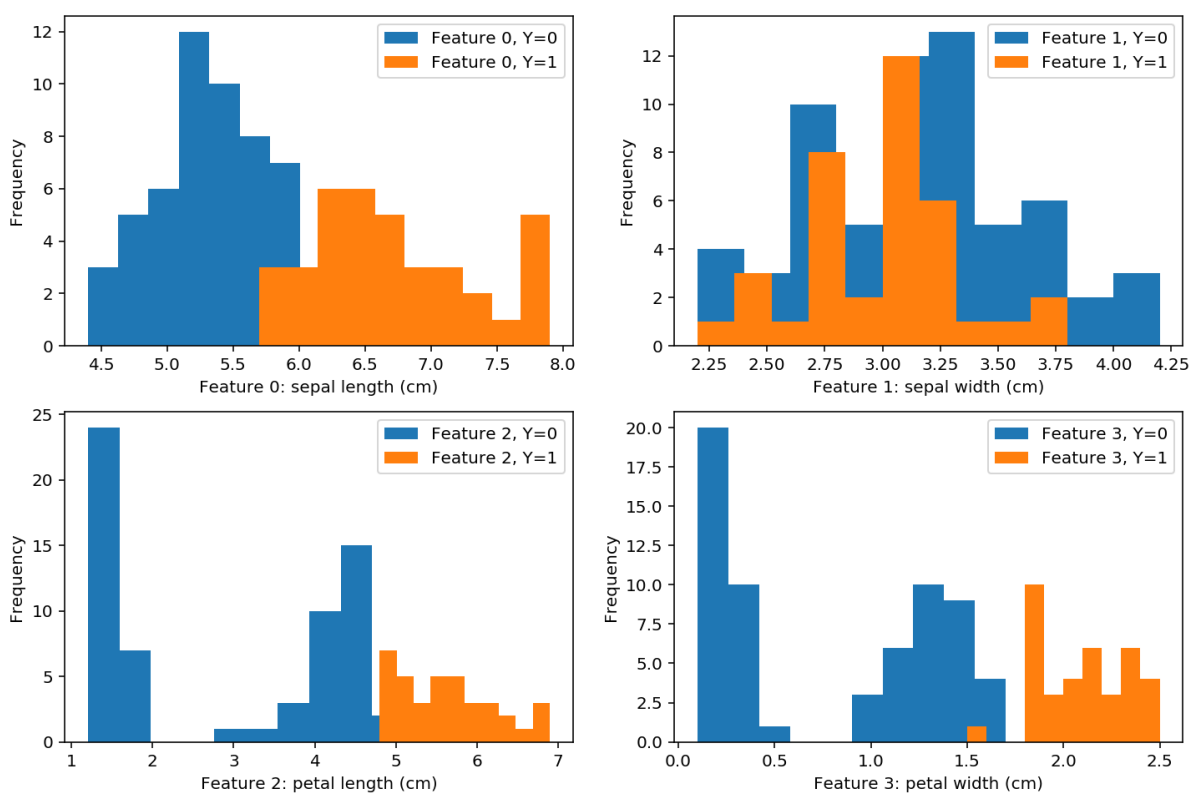
(50,)

```

In [48]: # Show the histograms of each feature.
plt.figure(figsize=(12,8))
for j in range(4):
    Xj_train = X_train[:,j]
    Xj_when_Y0_train = [Xj_train[i] for i in range(len(Xj_train)) if Y_train[i] == 0]
    Xj_when_Y1_train = [Xj_train[i] for i in range(len(Xj_train)) if Y_train[i] == 1]

    plt.subplot(2, 2, j+1)
    plt.hist(Xj_when_Y0_train, label='Feature {}, Y=0'.format(j))
    plt.hist(Xj_when_Y1_train, label='Feature {}, Y=1'.format(j))
    plt.xlabel('Feature {}: {}'.format(j, iris.feature_names[j]))
    plt.ylabel('Frequency')
    plt.legend()
plt.show()

```





```
In [49]: from sklearn.linear_model import LinearRegression
from sklearn.metrics import accuracy_score

reg = LinearRegression().fit(X_train, Y_train)
print(reg.coef_)
print(reg.intercept_)

w = reg.coef_
b = reg.intercept_

y_hat_train = np.dot(X_train, w) + b
label_train = (y_hat_train <= .5)
label_train = label_train*1

epilson_i_train = y_hat_train - Y_train
epilson_i_train_sqrtd = np.square(epilson_i_train)

y_hat_test = np.dot(X_test, w) + b
epilson_i_test = y_hat_test - Y_test
epilson_i_test_sqrtd = np.square(epilson_i_test)

[ 0.12975624  0.12249935 -0.11714156  0.67102651]
-1.1698768088050127
```

### Training error of the regression model.

```
In [50]: #Training error of the regression model
sum_epilson_i_train_sqrtd = np.sum(epilson_i_train_sqrtd)
regression_training_error = np.sqrt(sum_epilson_i_train_sqrtd/len(epilson_i_train_sqrtd))
print(regression_training_error)

0.27976412743241214
```

### Testing error of the regression model.

```
In [51]: #Testing error of the regression model.
sum_epilson_i_test_sqrtd = np.sum(epilson_i_test_sqrtd)
regression_test_error = np.sqrt(sum_epilson_i_test_sqrtd/len(epilson_i_test_sqrtd))
print(regression_test_error)

0.3310071344139558
```

### Training error of the classification model.

```
In [52]: #Training error of the classification model.
list = (label_train == Y_train)
list = 1*list
training_classification_error = sum(list) / len(list)
print(training_classification_error)
```

0.06

### Testing error of the classification model.

```
In [53]: #Testing error of the classification model.
label_test = (y_hat_test <= .5)
label_test = label_test*1
list_test = label_test == Y_test
test_classification_error = sum(list_test) / len(list_test)
print(test_classification_error)
```

0.14

## 6) Decision Stump

In this problem, we will perform a binary classification task on the Iris dataset. Again, this dataset has 150 data points, where each data point  $x \in \mathbb{R}^4$  has 4 features and its corresponding label  $y \in \{0, 1\}$ . To classify these 2 labels above, we decide to utilize a decision stump. The decision stump works as follows (for simplicity, we restrict our attention to uni-directional decision stumps):

• Given the  $j$ -th feature  $x_i(j)$  and a threshold  $Th_j$ , for data point  $i$ , the classification function is defined by  $y = f(x, j, Th_j)$  as:

$$f(x, j, Th_j) = \begin{cases} 1, & \text{if } x(j) > Th_j \\ 0, & \text{Otherwise} \end{cases}$$

Based on the decision stump above, we wish to write an algorithm to find the best feature and best threshold on training set to create a “best” decision stump, in a sense that such decision stump achieves the highest accuracy on training set.

Follow the instructions in the skeleton code and report:

```

In [36]: # Calculate the accuracy of prediction given feature, target and threshold.
def calc_acc(Xj, Y, thres):
    """
    Calculate the accuracy given feature, target and threshold.
    Xj:    j-th feature. This array only contains 1 feature for all
    data points,
           so the shape should be (count of data points,)
    Y:     Target array. Shape: (count of data points,)
    thres: Threshold.
    Return the accuracy of prediction.
    """

    # Step 1. Count the number of correct predictions and incorrect predictions.
    #           Here, for simplicity, we assume:
    #           If feature <= threshold, we predict it as Y = 0.
    #           If feature > threshold, we predict it as Y = 1.
    n_correct = 0
    n_incorrect = 0

    # ***** To be filled *****
    thres_arr = np.full((Xj.shape[0], ), thres)
    label = (Xj > thres_arr)
    label = 1*label

    correct = 1*(label == Y)
    n_correct = sum(correct)
    n_incorrect = len(Y) - n_correct

    # Step 2. Calculate the accuracy.
    acc = 1.0 * n_correct / (n_correct + n_incorrect)

    return acc

```

```

In [37]: # Show the histograms of each feature.
plt.figure(figsize=(12,9))

all_max_acc = 0.0 # Max training accuracy among all features.
all_thres = None # Threshold when reach the max training accuracy.
all_feature = None # Index of feature when reach the max training accuracy.

print(X_train)
# Loop over 4 features. j: index of current feature.
for j in range(4):
    # Get data.
    Xj_train = X_train[:,j] # Array of feature j.
    Xj_when_Y0_train = [Xj_train[i] for i in range(len(Xj_train)) if Y_train[i] == 0] # Array of feature j when Y = 0.
    Xj_when_Y1_train = [Xj_train[i] for i in range(len(Xj_train)) if Y_train[i] == 1] # Array of feature j when Y = 1.

    current_max_acc = 0.0 # Max training accuracy in current feature.
    current_thres = None # Threshold when reach the max accuracy in current feature.

    # Loop over all possible values for threshold. Here we consider 100 numbers between min and max of current feature.
    for thres in np.linspace(Xj_train.min(), Xj_train.max(), 100):
        # Calculate the accuracy on training data given feature, target and threshold.
        acc = calc_acc(Xj_train, Y_train, thres)
        # Update the current max accuracy if possible.
        if acc > current_max_acc:
            current_max_acc = acc
            current_thres = thres

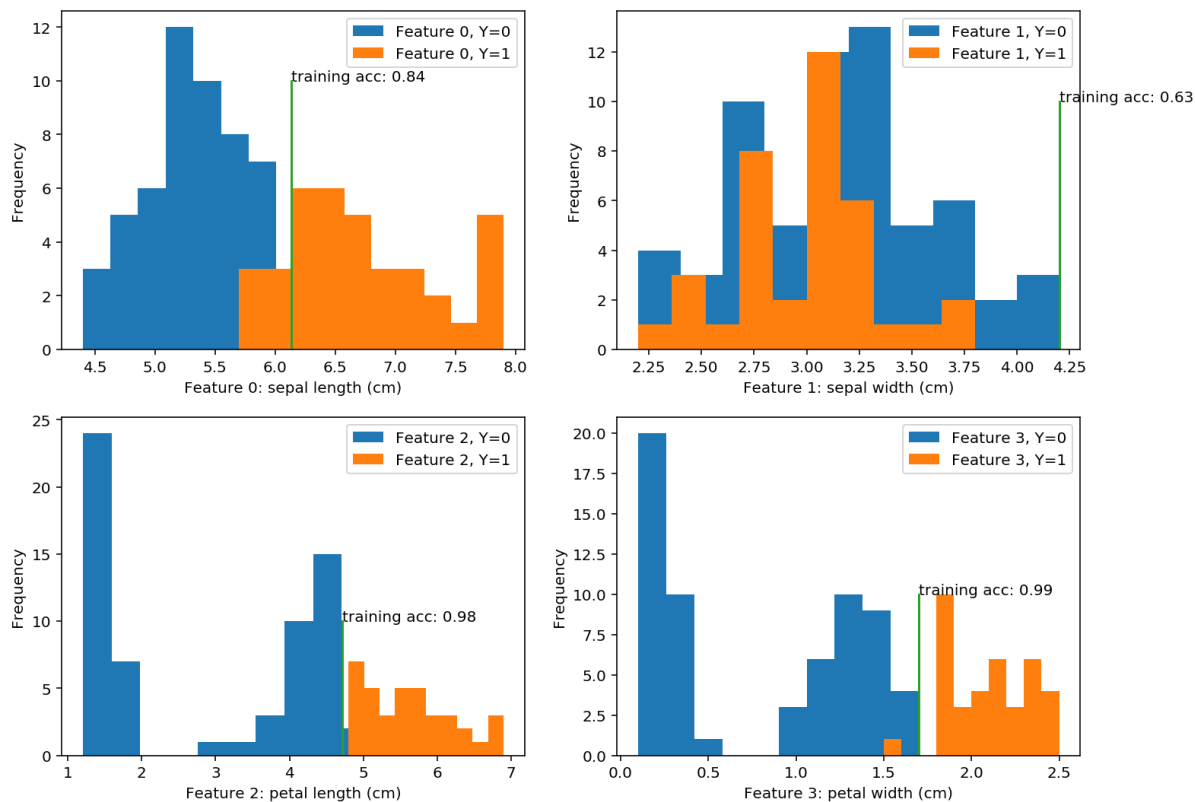
    # Update the max training accuracy among all features if possible.
    if current_max_acc > all_max_acc:
        all_max_acc = current_max_acc
        all_thres = current_thres
        all_feature = j

    # Plot the histograms and the best decision stump in current feature.
    plt.subplot(2, 2, j+1)
    plt.hist(Xj_when_Y0_train, label='Feature {}, Y=0'.format(j))
    plt.hist(Xj_when_Y1_train, label='Feature {}, Y=1'.format(j))
    plt.plot([current_thres, current_thres], [0, 10])
    plt.text(current_thres, 10, 'training acc: {}'.format(current_max_acc))
plt.xlabel('Feature {}: {}'.format(j, iris.feature_names[j]))
plt.ylabel('Frequency')
plt.legend()
plt.show()

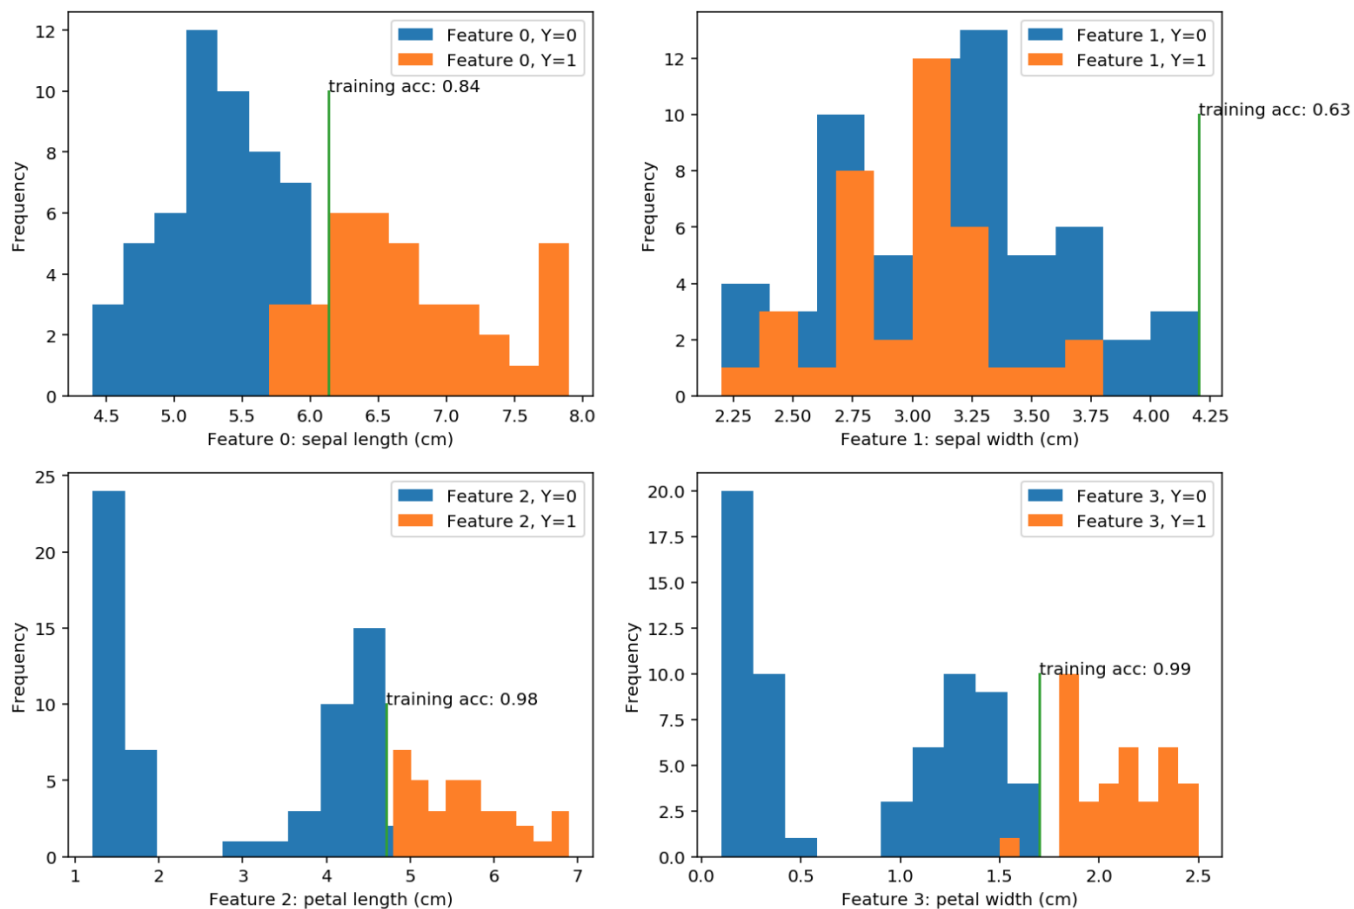
```

```
[ [5.8 4. 1.2 0.2]
[5.1 2.5 3. 1.1]
[6.6 3. 4.4 1.4]
[5.4 3.9 1.3 0.4]
[7.9 3.8 6.4 2. ]
[6.3 3.3 4.7 1.6]
[6.9 3.1 5.1 2.3]
[5.1 3.8 1.9 0.4]
[4.7 3.2 1.6 0.2]
[6.9 3.2 5.7 2.3]
[5.6 2.7 4.2 1.3]
[5.4 3.9 1.7 0.4]
[7.1 3. 5.9 2.1]
[6.4 3.2 4.5 1.5]
[6. 2.9 4.5 1.5]
[4.4 3.2 1.3 0.2]
[5.8 2.6 4. 1.2]
[5.6 3. 4.5 1.5]
[5.4 3.4 1.5 0.4]
[5. 3.2 1.2 0.2]
[5.5 2.6 4.4 1.2]
[5.4 3. 4.5 1.5]
[6.7 3. 5. 1.7]
[5. 3.5 1.3 0.3]
[7.2 3.2 6. 1.8]
[5.7 2.8 4.1 1.3]
[5.5 4.2 1.4 0.2]
[5.1 3.8 1.5 0.3]
[6.1 2.8 4.7 1.2]
[6.3 2.5 5. 1.9]
[6.1 3. 4.6 1.4]
[7.7 3. 6.1 2.3]
[5.6 2.5 3.9 1.1]
[6.4 2.8 5.6 2.1]
[5.8 2.8 5.1 2.4]
[5.3 3.7 1.5 0.2]
[5.5 2.3 4. 1.3]
[5.2 3.4 1.4 0.2]
[6.5 2.8 4.6 1.5]
[6.7 2.5 5.8 1.8]
[6.8 3. 5.5 2.1]
[5.1 3.5 1.4 0.3]
[6. 2.2 5. 1.5]
[6.3 2.9 5.6 1.8]
[6.6 2.9 4.6 1.3]
[7.7 2.6 6.9 2.3]
[5.7 3.8 1.7 0.3]
[5. 3.6 1.4 0.2]
[4.8 3. 1.4 0.3]
[5.2 2.7 3.9 1.4]
[5.1 3.4 1.5 0.2]
[5.5 3.5 1.3 0.2]
[7.7 3.8 6.7 2.2]
[6.9 3.1 5.4 2.1]
[7.3 2.9 6.3 1.8]
[6.4 2.8 5.6 2.2]
[6.2 2.8 4.8 1.8]
```

```
[6.  3.4 4.5 1.6]
[7.7 2.8 6.7 2. ]
[5.7 3.  4.2 1.2]
[4.8 3.4 1.6 0.2]
[5.7 2.5 5.  2. ]
[6.3 2.7 4.9 1.8]
[4.8 3.  1.4 0.1]
[4.7 3.2 1.3 0.2]
[6.5 3.  5.8 2.2]
[4.6 3.4 1.4 0.3]
[6.1 3.  4.9 1.8]
[6.5 3.2 5.1 2. ]
[6.7 3.1 4.4 1.4]
[5.7 2.8 4.5 1.3]
[6.7 3.3 5.7 2.5]
[6.  3.  4.8 1.8]
[5.1 3.8 1.6 0.2]
[6.  2.2 4.  1. ]
[6.4 2.9 4.3 1.3]
[6.5 3.  5.5 1.8]
[5.  2.3 3.3 1. ]
[6.3 3.3 6.  2.5]
[5.5 2.5 4.  1.3]
[5.4 3.7 1.5 0.2]
[4.9 3.1 1.5 0.1]
[5.2 4.1 1.5 0.1]
[6.7 3.3 5.7 2.1]
[4.4 3.  1.3 0.2]
[6.  2.7 5.1 1.6]
[6.4 2.7 5.3 1.9]
[5.9 3.  5.1 1.8]
[5.2 3.5 1.5 0.2]
[5.1 3.3 1.7 0.5]
[5.8 2.7 4.1 1. ]
[4.9 3.1 1.5 0.1]
[7.4 2.8 6.1 1.9]
[6.2 2.9 4.3 1.3]
[7.6 3.  6.6 2.1]
[6.7 3.  5.2 2.3]
[6.3 2.3 4.4 1.3]
[6.2 3.4 5.4 2.3]
[7.2 3.6 6.1 2.5]
[5.6 2.9 3.6 1.3]]
```



• All 4 histograms in last part of the code.



- The best feature, best threshold, training and test accuracy in last part of the code.

```
In [39]: # Use the best feature and best threshold on test set.  
  
Xj_test = X_test[:, all_feature] # Array of best feature.  
test_acc = calc_acc(Xj_test, Y_test, all_thres)  
print('Best feature: {}'.format(all_feature))  
print('Best threshold: {:.2f}'.format(all_thres))  
print('Training accuracy of best feature: {:.2f}'.format(all_max_acc))  
print('Test accuracy of best feature: {:.2f}'.format(test_acc))
```

```
Best feature: 3  
Best threshold: 1.70  
Training accuracy of best feature: 0.99  
Test accuracy of best feature: 0.90
```